



# KPLABS Course

Certified Kubernetes Security Specialist

Monitoring, Logging, and Runtime Security

ISSUED BY

Zeal

REPRESENTATIVE

[instructors@kplabs.in](mailto:instructors@kplabs.in)



# Module 1: Overview of Falco

## 1.1 Detection and Prevention

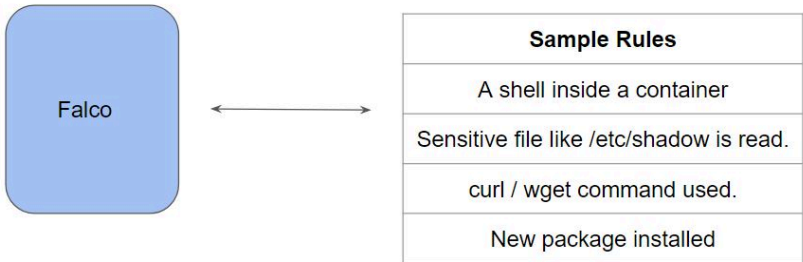
In Kubernetes, various features like Network Policies, RBAC, and others are primarily used for prevention purposes.

On the other hand, it is also important to add a certain set of detection measures that allows us to see what exactly is happening within the environment.



## 1.2 Basics of Falco

Falco is an open-source-based security tool that allows users to define a set of rules that will trigger an alert whenever the conditions are met.



### 1.3 Basic Rule Format

Falco rules are written in YAML and have a variety of required and optional keys.

rule	Name of the rule.
desc	Description of what the rule is filtering for.
condition	The logic statement that triggers a notification.
output	The message that will be shown in the notification.
priority	The "logging level" of the notification.

### 1.4 Sample Rule

Alerts whenever there is a shell spawned inside a container.


```
- rule: Detect bash in a container
  desc: You shouldn't have a shell run in a container
  condition: container.id != host and proc.name = bash
  output: Bash ran inside a container (user=%user.name command=%proc.cmdline %container.info)
  priority: INFO
```

## Module 2: Overview of Sysdig

### 2.1 Basics of Sysdig:

In a normal scenario of troubleshooting and performance monitoring, we make use of the following tools

Sysdig offers the functionality of these tools along with a lot more.



## 2.2 Interactive Options

Sysdig Utility comes with a command-line option (sysdig) as well as interface UI (csysdig)

Viewing: Processes For: whole machine  
Source: Live System Filter: evt.type=switch

PID	CPU	USER	TH	VIRT	RES	FILE	NET	Command
123827	6.50	root	8	577M	159M	0	0.00	csysdig
102296	6.00	root	9	1G	334M	0	15.22K	kube-apiserver --advertise-address=172.31.59.221
471	3.00	root	18	2G	107M	0	3.39K	/usr/bin/kubelet --bootstrap-kubeconfig=/etc/kube
102255	2.50	root	7	793M	100M	0	17.89K	kube-controller-manager --allocate-node-cidrs=tru
663	1.50	root	27	1G	114M	0	0.00	/usr/bin/dockerd -H fd:// --containerd=/run/conta
1966	1.50	root	14	2G	73M	70K	2.98K	etcd --advertise-client-urls=https://172.31.59.22
4714	0.50	root	9	730M	38M	0	893.50	/coredns -conf /etc/coredns/CoreFile
491	0.50	root	33	1G	49M	0	0.00	/usr/bin/containerd
3550	0.00	root	9	1G	36M	0	0.00	/opt/bin/flanneld --ip-masq --kube-subnet-mgr
123819	0.00	ubuntu	1	6M	4M	0	0.00	/usr/lib/openssh/sftp-server
1689	0.00	root	10	106M	6M	0	0.00	containerd-shim -namespace moby -workdir /var/lib
488	0.00	root	1	17M	8M	0	0.00	/lib/systemd/systemd-logind
4368	0.00	root	10	106M	5M	0	0.00	containerd-shim -namespace moby -workdir /var/lib
4353	0.00	root	10	106M	5M	0	0.00	containerd-shim -namespace moby -workdir /var/lib
102256	0.00	root	9	729M	46M	0	1.51K	kube-scheduler --authentication-kubeconfig=/etc/k
534	0.00	www-data	1	57M	5M	0	0.00	nginx: worker process
3265	0.00	root	10	106M	6M	0	0.00	containerd-shim -namespace moby -workdir /var/lib
512	0.00	root	1	6M	2M	0	0.00	/sbin/agetty -o -p -- -u --noclear tty1 linux
1641	0.00	root	11	853M	28M	0	0.00	/snap/amazon-ssm-agent/2996/ssm-agent-worker
3123	0.00	root	1	964K	4K	0	0.00	pause
458	0.00	root	1	8M	3M	0	0.00	/usr/sbin/cron -f

## 2.3 Running sysdig

In its simplest form, when you run sysdig, you will see all the system calls that are happening within the system.

```
52 07:50:42.829676193 1 <NA> (0) > switch next=1157782 pgft_maj=0 pgft_min=0 vm_size=0 vm_rss=0 vm_swap=0
54 07:50:42.829686696 0 sysdig (1171310) > switch next=1155766 pgft_maj=0 pgft_min=2183 vm_size=447872 vm_rss=25612
vm_swap=0
55 07:50:42.829687819 1 <NA> (1157782) > switch next=1170533(sshd) pgft_maj=0 pgft_min=0 vm_size=0 vm_rss=0 vm_swap
=0
56 07:50:42.829690356 0 <NA> (1155766) > switch next=1171310(sysdig) pgft_maj=0 pgft_min=0 vm_size=0 vm_rss=0 vm_sw
ap=0
57 07:50:42.829696854 1 sshd (1170533) < select res=1
58 07:50:42.829697168 0 sysdig (1171310) > switch next=1155766 pgft_maj=0 pgft_min=2183 vm_size=447872 vm_rss=25612
vm_swap=0
59 07:50:42.829699893 0 <NA> (1155766) > switch next=1171310(sysdig) pgft_maj=0 pgft_min=0 vm_size=0 vm_rss=0 vm_sw
ap=0
60 07:50:42.829702032 1 sshd (1170533) > rt_sigprocmask
61 07:50:42.829703224 1 sshd (1170533) < rt_sigprocmask
62 07:50:42.829704021 1 sshd (1170533) > rt_sigprocmask
63 07:50:42.829704472 1 sshd (1170533) < rt_sigprocmask
64 07:50:42.829707634 1 sshd (1170533) < clock_gettime
65 07:50:42.829708525 1 sshd (1170533) < clock_gettime
66 07:50:42.829710931 1 sshd (1170533) > read fd=15(<f>/dev/ptmx) size=16384
67 07:50:42.829714826 1 sshd (1170533) < read res=1201 data=1 07:50:42.736703000 0 container:0d65633084bf (-1) > co
ntainer_json={"container":
```

## 2.4 Filters

Since you will get a huge amount of data when monitoring system calls, you can use sysdig with filters to make the output more fine-grained.

```
root@ip-172-31-59-221:~# sysdig proc.name=nano
310161 07:55:39.213391058 0 nano (1173625) < execve res=0 exe=nano args=test.txt. tid=1173625(nano) pid=1173625(nan
o) ptid=1173599(bash) cwd= fdlimit=1024 pgft_maj=1 pgft_min=24 vm_size=652 vm_rss=4 vm_swap=0 comm=nano cgroups=cpu
set=/cpu=/user.slice.cpuacct=/user.slice.io=/user.slice.memory=/user.slice... env=SHLL=/bin/bash.SUDO_GID=1000.SUD
O_COMMAND=/usr/bin/su.SUDO_USER=ubuntu.PWD=/... tty=34822 pgid=1173625(nano) loginuid=1000
310162 07:55:39.213423440 0 nano (1173625) > brk addr=0
310163 07:55:39.213424690 0 nano (1173625) < brk res=563F7341E000 vm_size=652 vm_rss=4 vm_swap=0
310164 07:55:39.213434121 0 nano (1173625) > arch_prctl
310165 07:55:39.213434779 0 nano (1173625) < arch_prctl
310166 07:55:39.213459275 0 nano (1173625) > access mode=4(R_OK)
310167 07:55:39.213464446 0 nano (1173625) < access res=-2(ENOENT) name=/etc/ld.so.preload
310168 07:55:39.213469242 0 nano (1173625) > openat
310169 07:55:39.213476122 0 nano (1173625) > fstat fd=3
310170 07:55:39.213477566 0 nano (1173625) < fstat res=0
```

## 2.5 Sysdig Chisels

Sysdig's chisels are little scripts that analyze the sysdig event stream to perform useful actions

```
root@ip-172-31-59-221:~# sysdig -cl

Category: Application
-----
httplog           HTTP requests log
httpstop          Top HTTP requests
memcachelog       memcached requests log

Category: CPU Usage
-----
spectrogram       Visualize OS latency in real time.
subsecoffset      Visualize subsecond offset execution time.
topcontainers_cpu Top containers by CPU usage
topprocs_cpu      Top processes by CPU usage

Category: Errors
-----
topcontainers_error Top containers by number of errors
topfiles_errors    Top files by number of errors
topprocs_errors    Top processes by number of errors
```

## Module 3: Writing Falco Rules

### 3.1 Basic Rule Format

Falco rules are written in YAML and have a variety of required and optional keys.

rule	Name of the rule.
desc	Description of what the rule is filtering for.
condition	The logic statement that triggers a notification.
output	The message that will be shown in the notification.
priority	The “logging level” of the notification.

Here is a sample Falco rule:

```
- rule: Detect bash in a container
  desc: You shouldn't have a shell run in a container
  condition: container.id != host and proc.name = bash
  output: Bash ran inside a container (user=%user.name command=%proc.cmdline %container.info)
  priority: INFO
```

3.2 Rule Elements

A Falco rules file is a YAML file containing three types of elements:

Elements	Description
Rules	Conditions under which an alert should be generated. A rule is accompanied by a descriptive output string that is sent with the alert.
Macros	Macros provide a way to name common patterns and factor out redundancies in rules.
Lists	Collections of items that can be included in rules, macros, or other lists

Module 4: Falco Rule Writing - Exam Perspective

XYZ use-case

Format as follows

time	uid	process-name
------	-----	--------------

Capture log for 25 seconds and store it in /tmp/log.txt